


Wireguard FPGA



Advanced co-simulation
verification environment



Simon Southwell aka Wyvern Semiconductors



Freelance system architect, logic development and software consultant, specializing in co-simulation, HPC and network solutions. Generating and collaborating on open-source projects, mentoring and presenting. Contributor to the OSVVM project, adding and supporting co-simulation capabilities.

simon@anita-simulators.org.uk

[LinkedIn: Simon Southwell](#)

Cambridge, United Kingdom

Focus of Today's Presentation



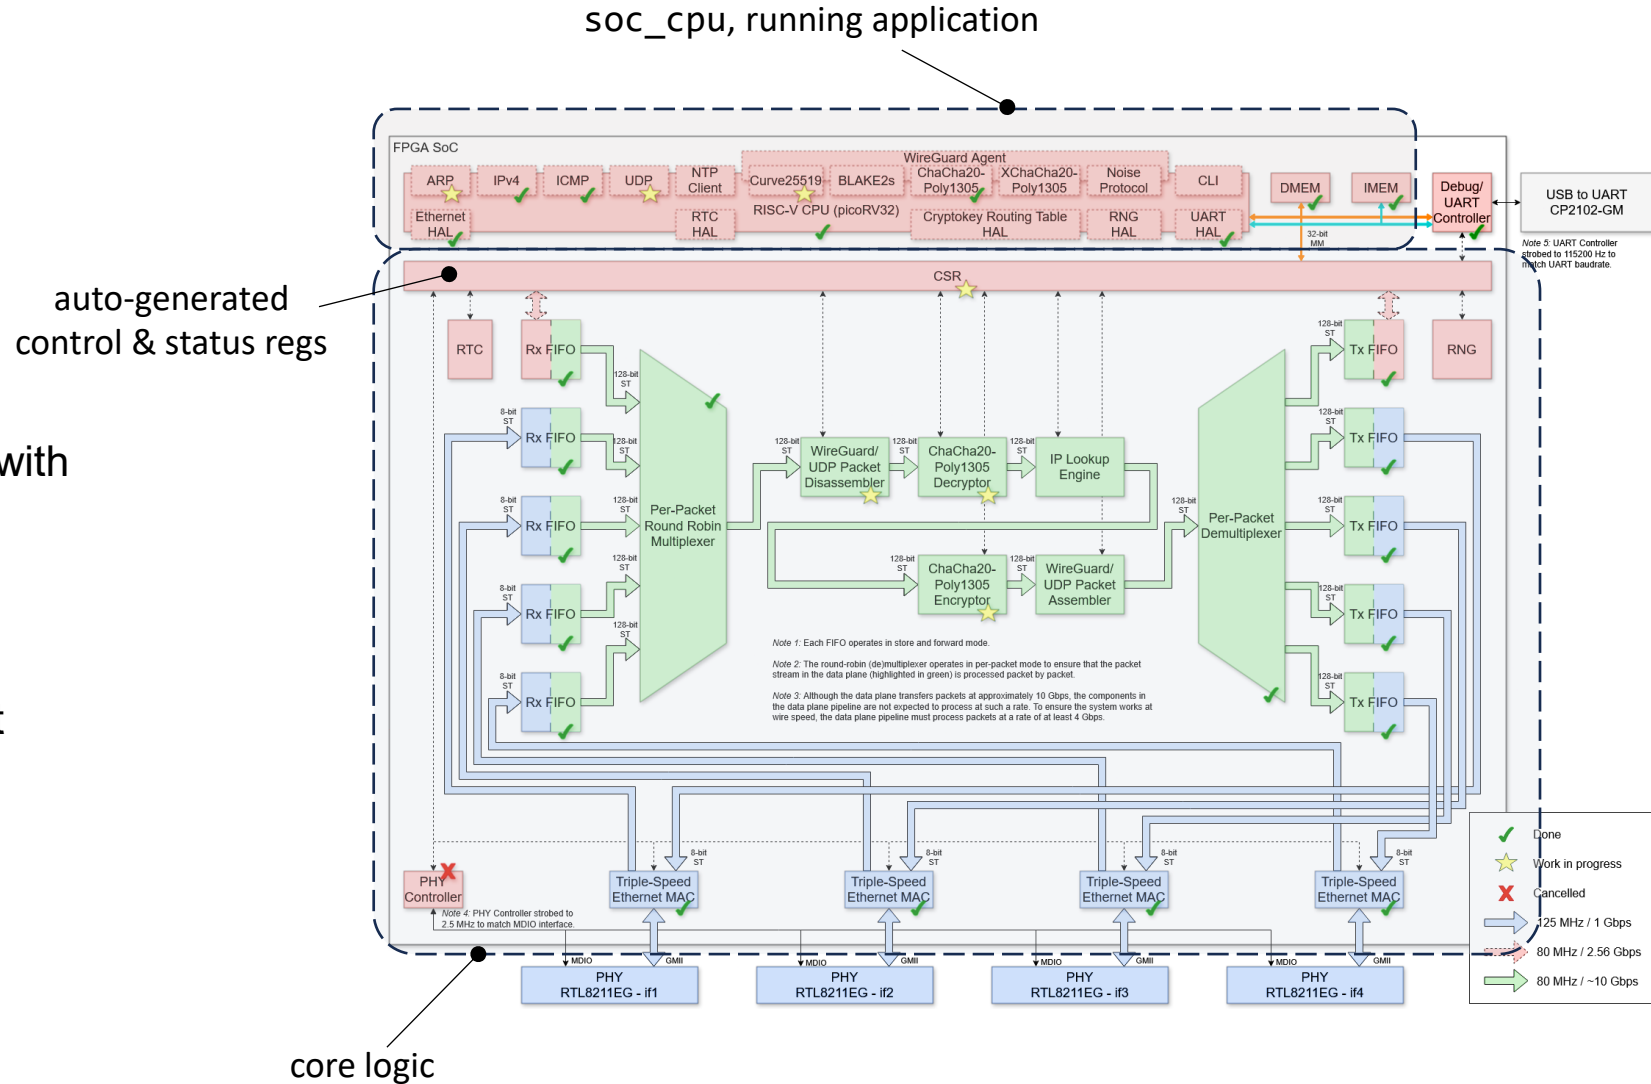
- The main thrust of the presentation will be on the co-simulation used in the Wireguard FPGA project's top level logic simulation
- Will give a very brief overview of the Wireguard FPGA project for context
- A summary of the overall verification architecture
- A look at the open-source individual components to add advance co-simulation capabilities
 - A "Virtual Processor" co-simulation element
 - A sparse memory model, in C, integrated into the simulation
 - A RISC-V Instruction Set Simulator, integrated into the Virtual Processor
 - Auto-generation of Hardware Abstraction Layer code from SystemRDL Description
 - An Ethernet UDP/IPv4 C++ model, running on a Virtual Processor to drive the Ethernet Ports
- How all these are combined to give a cohesive top level test bench for co-simulation and co-development



The Wireguard FPGA in a Nutshell



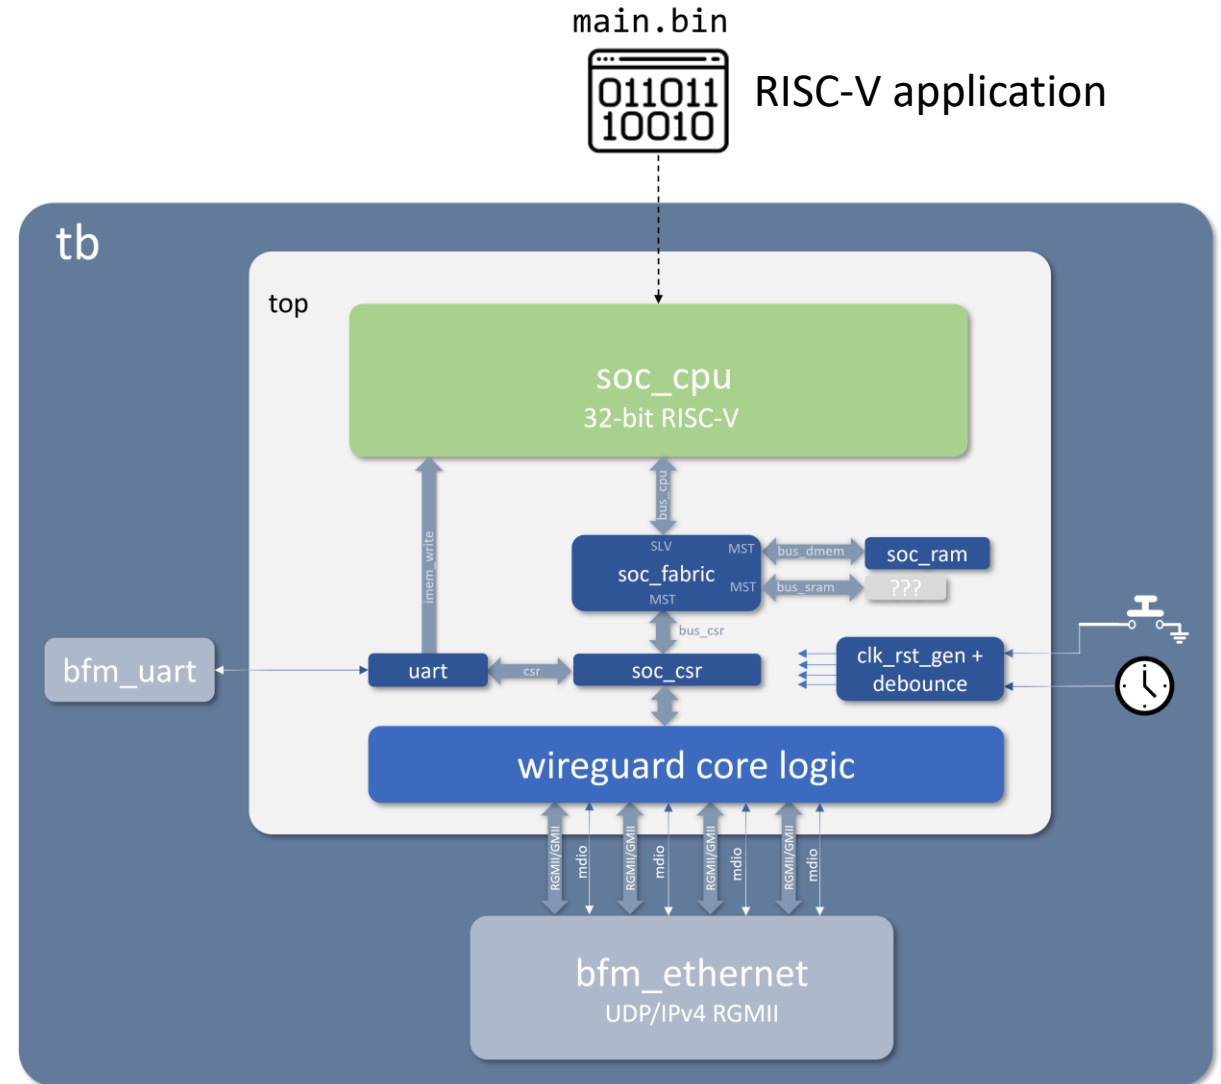
- Implemented by the young engineers at [Chili.CHIPS*ba](https://chili.chips.ba) out of Bosnia and with support from the [NLnet Foundation](https://nl.net/foundation)
- An open-source SystemVerilog logic implementation of the open-source Wireguard protocol for creating VPNs
- Targets inexpensive hardware platform with four 1000Base-T ports
- Using a commodity Artix7 FPGA
- Done in a self-sufficient way, i.e. without requiring a PC host
- Supported by open-source tools
 - E.g. [PipelineC](https://github.com/pipelineC) (Julian Kemmerer)
- All gateway written in standard Verilog/SystemVerilog



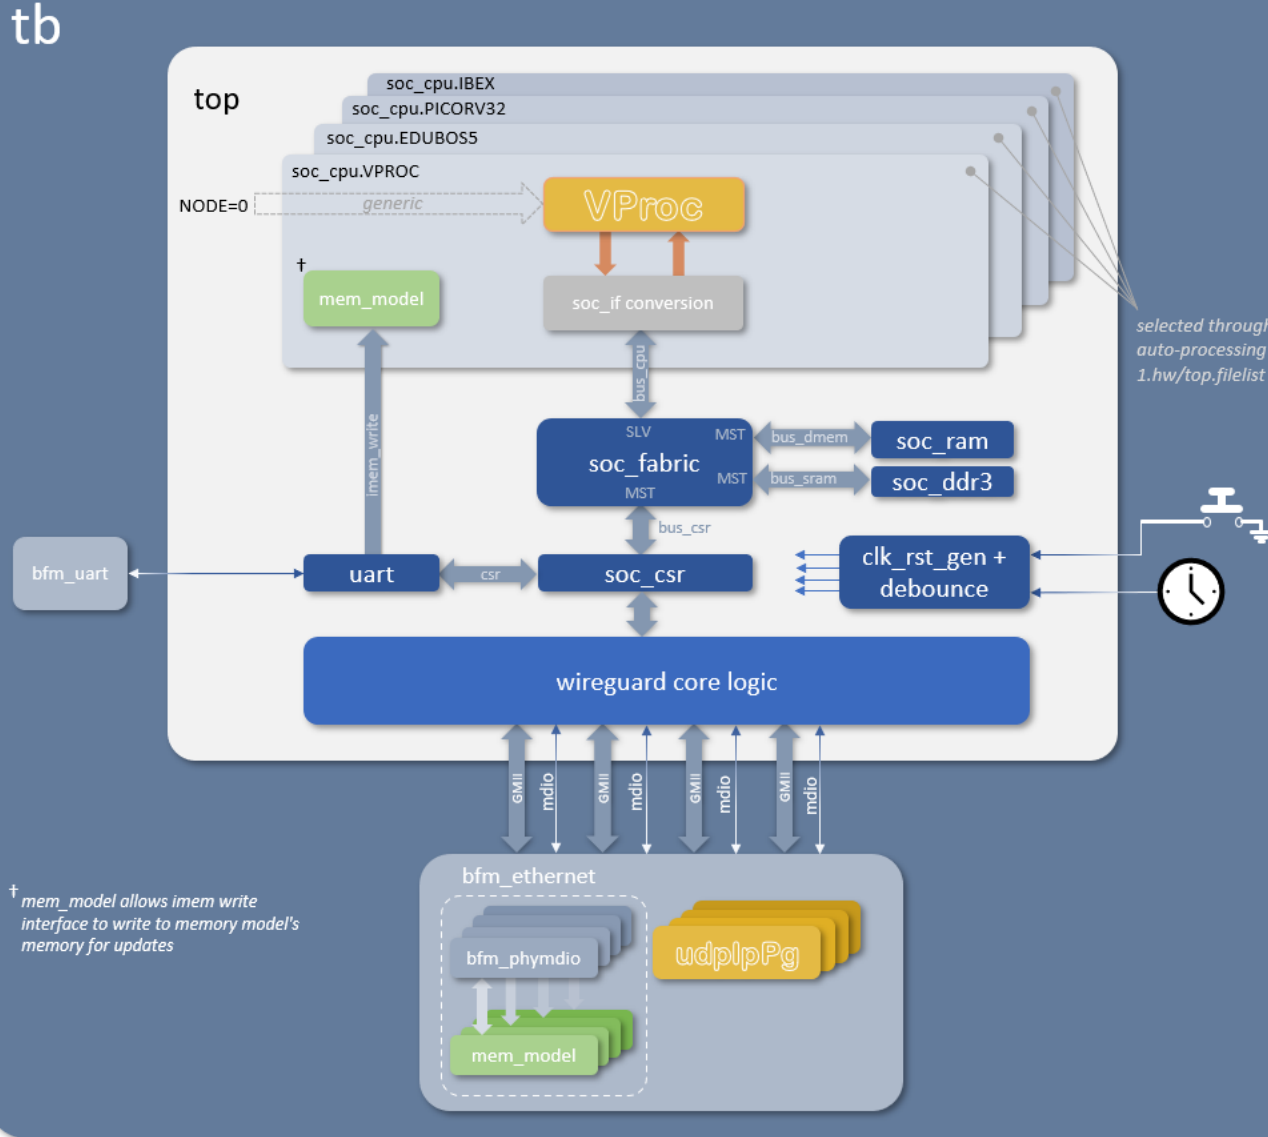
Wireguard Basic Top Level Test Bench



- Fairly standard top level logic simulation test bench
 - The illustration shows a simplified block diagram
- For the WG project, [Verilator](#) (Wilson Snyder et. al.) and SystemVerilog are chosen, with [GTKWave](#) (Anthony Bell) or [Surfer](#) (Linköping University) for displaying waveforms
- Wireguard's external Interfaces are connected to driver modules
 - Ethernet model with UDP/IPv4 over GMII and MDIO
 - UART that can update code in memory and has some debugging features
- The IP itself has a 32-bit RISC-V processor (soc_cpu) with RTL based around the open-source [PicoRV32](#) RISC-V core (Clifford Wolf et. al), though could use others
- RISC-V test code or application software can be loaded and run on the processor
- But there lies some **hidden secrets** that extend this generic test bench structure for true co-simulation!



Wireguard VProc Based Top Level Test Bench



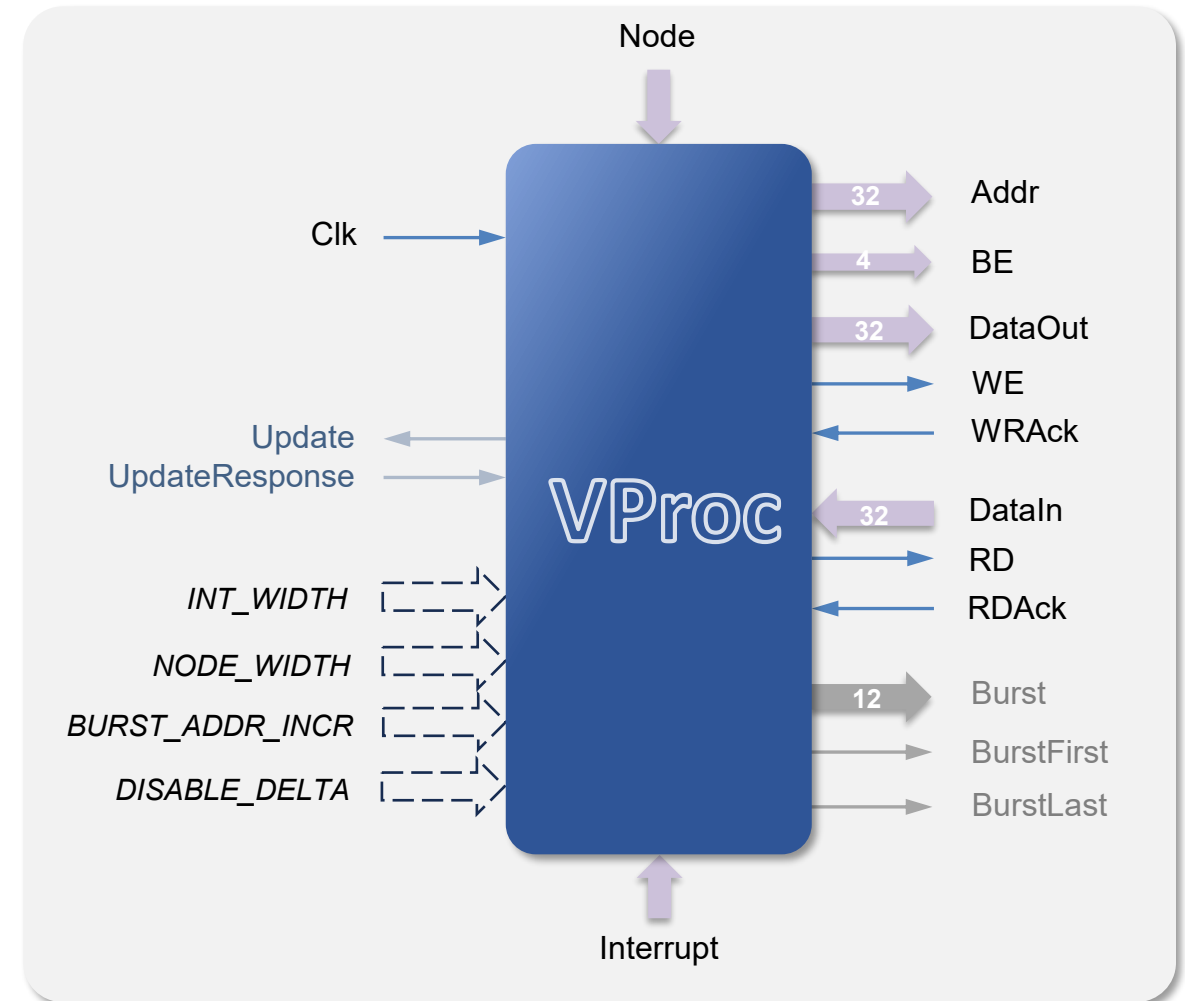
- Means are provided which can allow selection between flavours of soc_cpu components between various open source RISC-V RTL cores within the Wireguard FPGA IP
 - examples: PicoRV32, [IBEX](#) (lowRISC) and [EDUBOS](#) (Tarik Ibrahimovic)
- One selection, though, chooses a "Virtual Processor" known as [VProc](#) (wyvernSemi)
- Memory is modelled, in C, and "ports" to the model are made available in the test bench as HDL components via "mem_model" (wyvernSemi)
 - Programs running on the virtual processor can access the memory C model directly via its API
- Ethernet is driven by a UDP/IPv4 pattern generator C++ model, based around *VProc—udpPglp* (wyvernSemi)
- MDIO is driven by HDL but has access to the memory model
- This, then, is our route to true co-simulation and co-development, so let's look at these co-sim components

What is the *VProc* virtual processor?



- Can "run" a normal user C or C++ program, compiled for the host machine (PC or Linux Workstation), which can drive an address bus on an HDL component instantiated in a logic simulation
- The *VProc* HDL 'processor' Component has a generic memory mapped master interface
 - For WG, wrapped in an `soc_cpu` BFM for specific '`soc_if`' bus protocol
 - Optional Burst ports. Not used for Wireguard.
 - Supports interrupts. Not used for Wireguard
 - Has 'delta cycle' update capability (more later)
- In WG, DPI-C is used to connect between HDL and C/C++
 - Other simulators and programming interfaces supported
 - VHDL also supported
- Provides a C or C++ API to drive address bus
- Can instantiate multiple *VProcs*
 - Called 'nodes'.
 - Each must have a unique node number on **Node** port
- Each node's software has a specific "main" entry point for user code
 - E.g. `VUserMain0()` for node 0
 - c.f. `WinMain` for Windows graphics programs

The *VProc* HDL component



VProc's Software API



- User writes normal C or C++ code that is compiled into a static library and linked to the *Verilator* executable
 - Other simulators use a shared object loaded at run-time
- The basic C and C++ APIs:
 - Low Level C API, e.g.

```
VWrite (addr, data, delta, node);
VWriteBE(addr, data, be, delta, node);
VRead (addr, *data, delta, node);
VTick (ticks, node);
```
 - C++ *VProc* class, e.g.

```
vp0->writeByte (addr, data, delta=0);
vp0->writeHword (addr, data, delta=0);
vp0->writeWord (addr, data, delta=0);
vp0->readByte (addr, *data, delta=0);
vp0->readHword (addr, *data, delta=0);
vp0->readWord (addr, *data, delta=0);
vp0->tick (ticks);
```
- Using the API directly is suitable for writing custom tests to drive bus of `soc_cpu` and reading and writing registers or memory etc.
 - But we can do better than that!

Simple use of VProc's C++ API

```
#include "VProcClass.h"

static const int node = 0;

extern "C" void VUserMain0(void) // VProc "main" entry for node 0
{
    // API constructor
    VProc* vp0 = new VProc(node); // VProc API object for node 0

    vp0->tick(100); // Wait a bit

    uint32_t addr = 0x10001000; // Location in DMEM
    uint32_t wdata = 0x900dc0de;

    vp0->writeWord(addr, wdata);
    VPrint("Written 0x%08x to addr 0x%08x\n", wdata, addr);

    vp0->tick(10); // Emulate some processing (10 clock cycles)

    uint32_t rdata;
    vp0->readWord(addr, &rdata);

    if (rdata == wdata)
        VPrint("Read back 0x%08x from addr 0x%08x\n", rdata, addr);
    else
        VPrint("***ERROR: data mismatch at addr = 0x%08x\n", addr);

    // Sleep forever (and allow simulation to continue)
    while(true)
        vp0->tick(GO_TO_SLEEP);
}
```



Sparse Memory Model (*mem_model*)



- A sparse memory model in C is connected to an HDL component: [mem_model](#)
- *VProc* user code can access memory directly using the memory model's C API (`mem.h`). E.g.

```
WriteRamByte (addr, data, memnode)
WriteRamHword (addr, data, le, memnode)
WriteRamWord (addr, data, le, memnode)
ReadRamByte (addr, memnode)
ReadRamHword (addr, le, memnode)
ReadRamWord (addr, le, memnode)
```

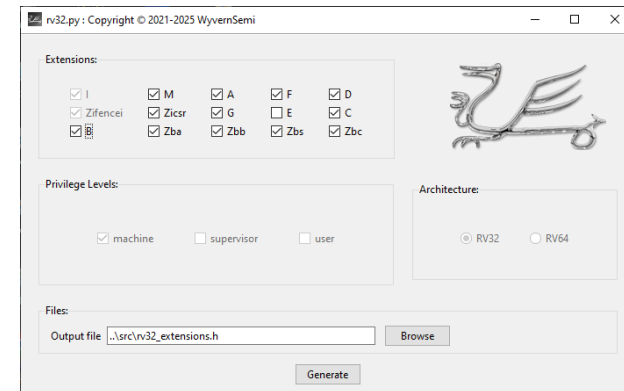
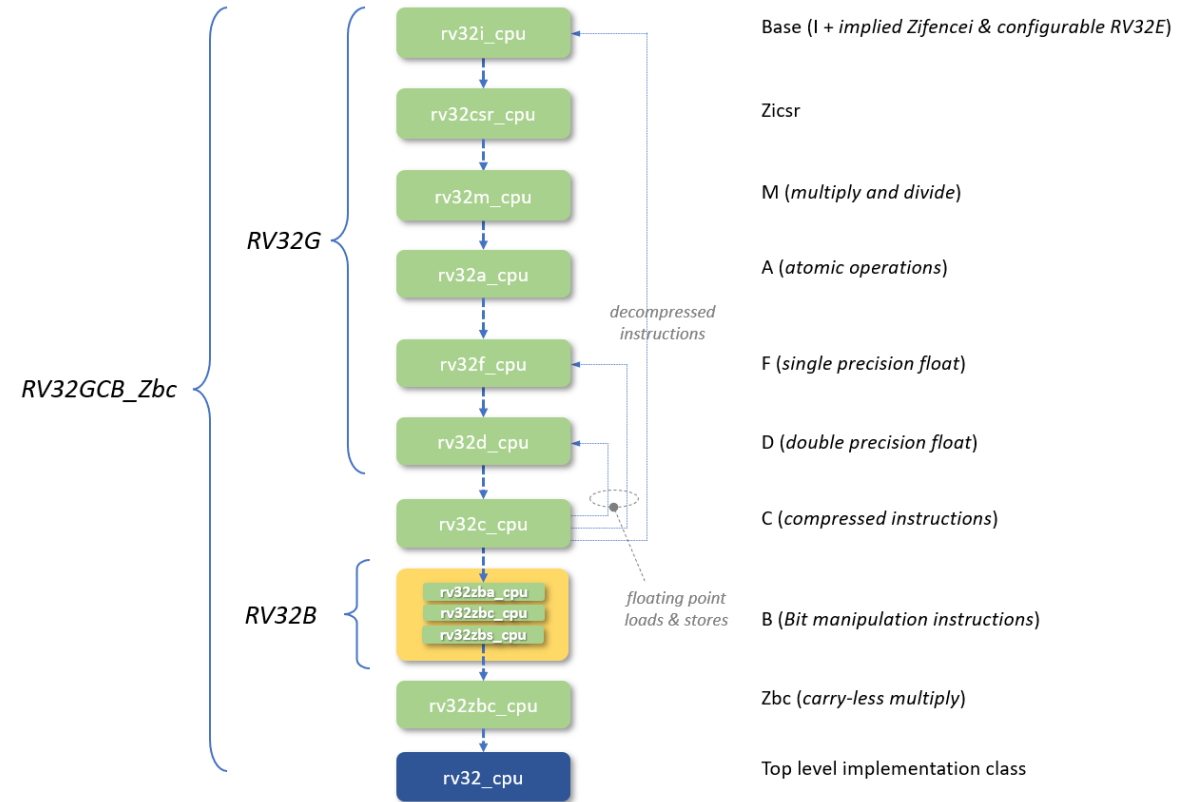
- The sparse memory model in C can implement up to a 2^{64} byte address space
- Connected to HDL with DPI for *Verilator* in this project
 - Gives access to C memory model from HDL
 - Module has various Altera Avalon style ports
- Can instantiate multiple `mem_model` components
 - NB: Each is a 'port' to the same address space (`memnode = 0`), *not* a new memory



The *rv32* RISC-V Instruction Set Simulator




- One of the "Programs" that can be run on the Virtual Processor is a C++ model of a RISC-V core
- *rv32* (wyvernSemi) is a C++ instruction set simulator model (ISS) for RISC-V, configurable up to RV32GCB_E_Zicsr_Zbc specifications
 - Can be configured to implement subsets of this (min RV32I)
 - Each RISC-V extension has own class and inherits the derived class before it to build up to a certain specification
- Can be compiled as a standalone executable or as a library for integration with other code (as done for WG)
- Can register external callback function, called whenever a memory access is made. Callback can decide to processor or hand back.
- Has GDB remote debug interface
- Has internal instruction timing model that can be configured to model different RISC-V cores' instruction timings
 - E.g. the PicoRV32 used in Wireguard FPGA
- On Wireguard, the ISS can be selected as the user code to run on the **soc_cpu.VPROC** component
 - A simple software layer is required to integrate with simulation (more later)
 - The RISC-V compiled application code is then run on the ISS



HAL Auto-generation for Co-simulation

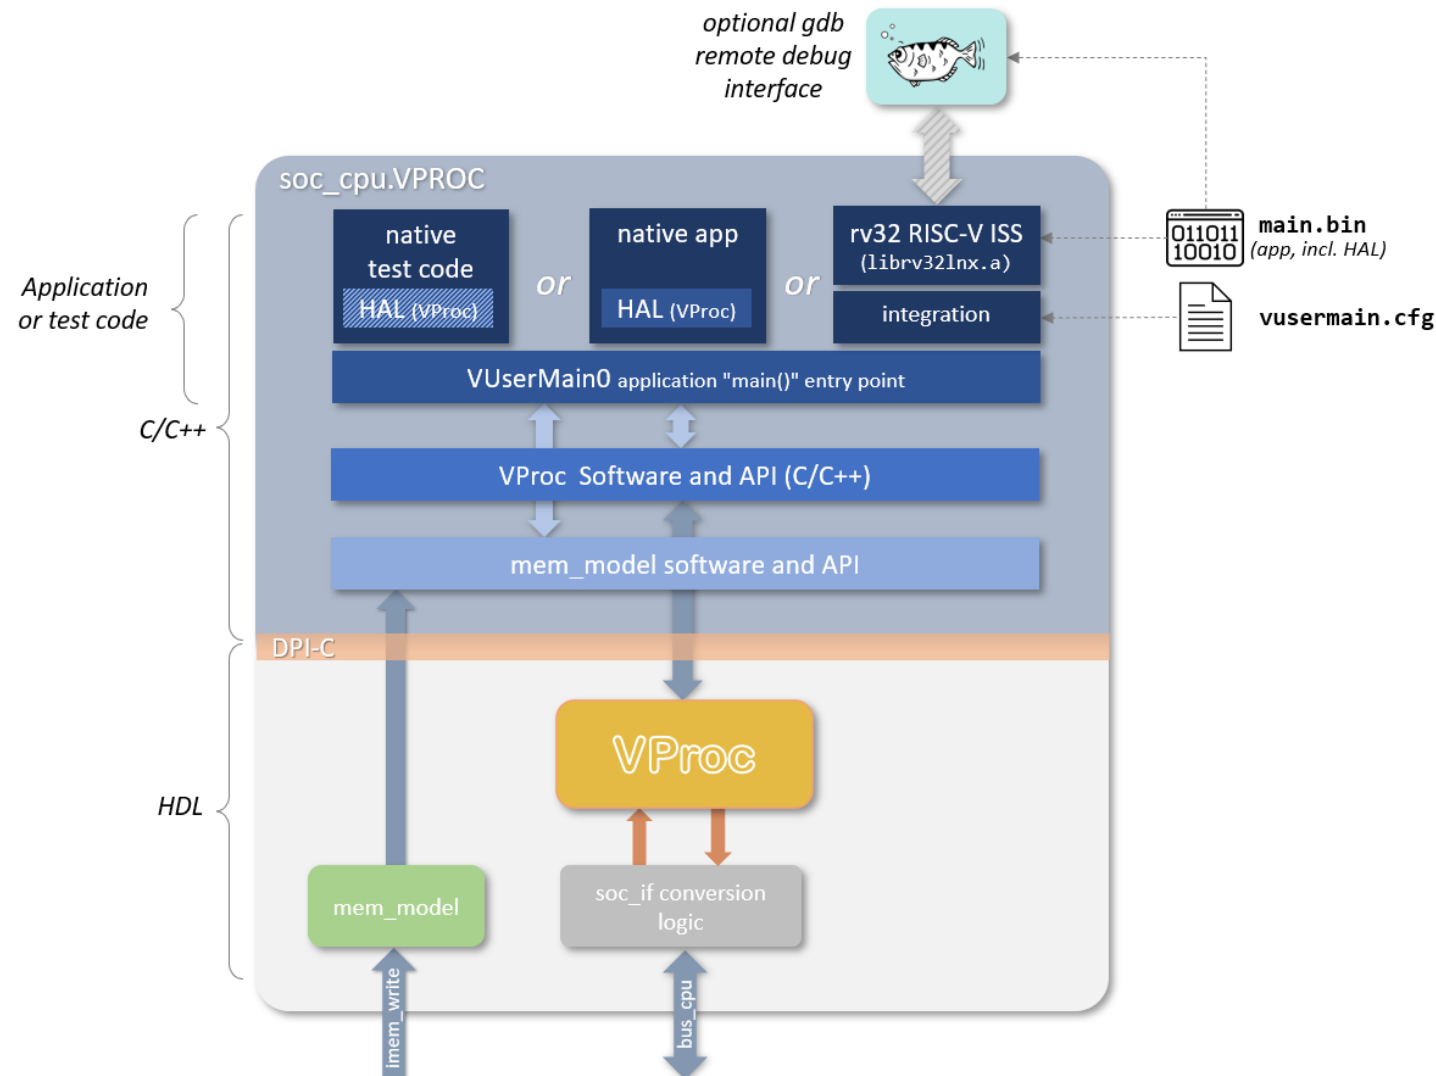


- Wireguard uses a SystemRDL (an [Accellera standard](#)) description of its CSR registers
- SystemRDL description parsed by the open-source [peakrdl](#) tools (Alex Mykyta et. al) :
 - Used to generate CSR RTL
 - Used to generate S/W hardware abstraction layer (HAL)
- Would like to compile the target application natively for host computer
 - Present the same HAL API to both native and target compiled code
 - Target (RISC-V) HAL calls *peakrdl* output
 - Native HAL calls *VProc* API functions
 - *systemrdl-compiler*, bundled with *peakrdl* tools, is used to generate the two types of co-simulation HAL
- With this, the Wireguard application can be developed in tandem with RTL without the need to model RISC-V
 - Will run faster
 - Functionally the same
 - Timings, though, will be approximate but good for fast functional verification
- Why not use the *peakrdl* output for both?
 - Use of bitfields negates overloading techniques previously used on other HALs



PeakRDL-regblock
PeakRDL-cheader
systemrdl-compiler

Putting it All Together for soc_cpu.VProc

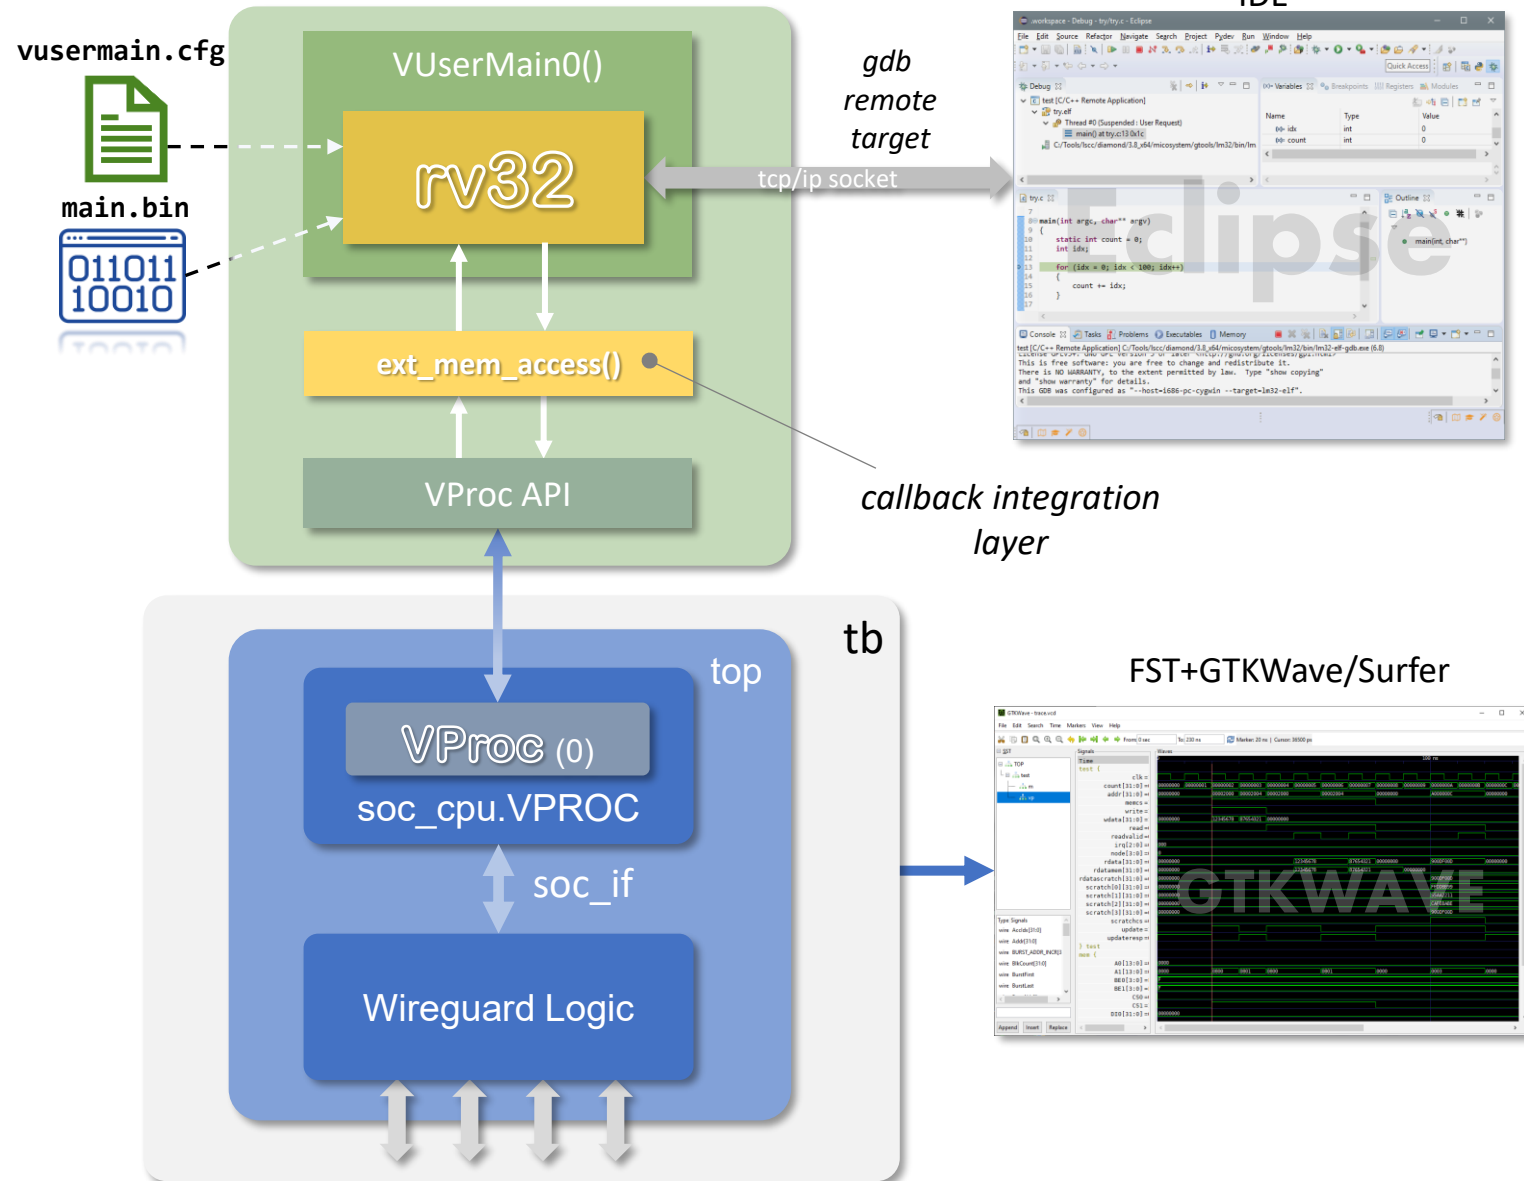


- Three use cases available for the TB:
 - Native test code to drive `soc_cpu` bus using `VProc` API directly, as discussed previously (or HAL)
 - Application software compiled natively with modified VProc HAL (same API as code compiled for RISC-V target)
 - `rv32` RISC-V ISS and integration code to run application for target.
- Debugging code a non-negotiable requirement
 - Native code can be debugged with normal gdb for host programs
 - RISC-V code can be debugged with remote gdb features of `rv32` ISS
- The `rv32` ISS can be configured at run time via a `vusermain.cfg` file
- Uses internal sparse memory model
 - `mem_model` HDL components give access from logic
 - `VProc` code can use it's direct C API
 - Thus software and hardware share the same memory space

Integrating *rv32* ISS into Wireguard TB



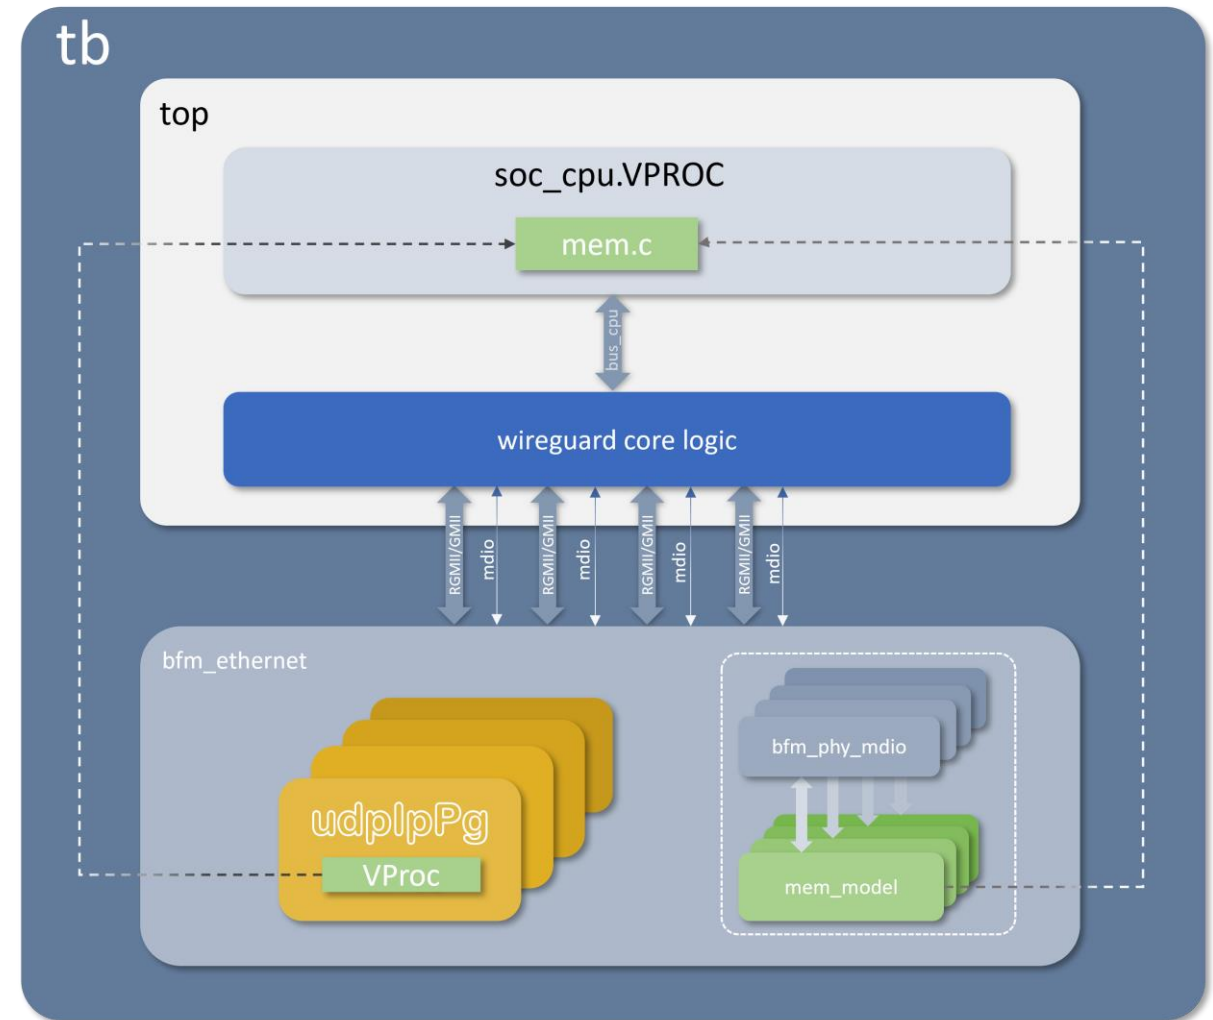
- *rv32* has a memory access callback feature
 - All loads/stores are 'offered' to a registered user callback function, along with the model's reckoning of time
 - Callback makes appropriate *VProc* read and write calls and to '*VTick*' to synchronise simulator with CPU's sense of time
 - If callback doesn't process access, returns access back to *rv32* to process
- Wireguard TB co-sim integration layer is implemented in this callback
 - If access within a configured range, makes accesses to CSR registers via *VProc* API
 - Otherwise uses *mem_model* C API to store in common address space



Ethernet UDP/IPv4 Driver VIP



- Based on [udplpPg](#) model (wyvernSemi) built around *VProc*
 - Developed from the [tcplpPg](#) model (wyvernSemi)
- Consists of a C++ model to generate and decode UDP/IPv4 ethernet packets for Gigabit Ethernet and run on a *VProc* component.
- *VProc* can be used to memory map ports on a module and drive multiple signals using its delta cycle feature
 - Any arbitrary interface can thus be driven
 - A small layer provides an API to the component ports
- The HDL components have GMII interfaces
 - An RGMII converter module is also provided
- An MDIO BFM HDL module is also available
 - Uses *mem_model* component to read from and write to (when receiving register accesses) configured segments of memory
- The *VProc* software of *udplpPg* and the *mem_model* used by the MDIO BFM both have access to the same memory space that *soc_cpu.VPROC* has
 - The test programs on *soc_cpu.VPROC* can now set patterns used by the other models or read data written by them, closing the loop for end-to-end verification



Conclusions



- From a starting point of a "classic" test bench architecture, have carefully deployed various readily available open source tools to make this into a true co-simulation and co-development verification environment
- A "Virtual Processor" provided the means to seamlessly bridge the HDL to C/C++ gap to "run" a natively compiled user program on an HDL processor core in a logic simulation
- A C memory model co-sim component gave a common memory space between HDL and user software
- A RISC-V instruction set simulator was integrated with VProc to allow the target's application code to be run in simulation as well as on the platform
- An auto-generated HAL, with a common API for target and VProc allowed the application to be compiled and run natively
- A soft UDP/IPv4 model was deployed to drive the ethernet ports, reutilising the virtual processor
- All this allows various options to be utilised, trading speed of execution, versus accuracy of timing
 - RTL processor (PicoRV32) with full timing accuracy but slow
 - rv32 ISS, high timing accuracy, but fast
 - Natively compiled application, lower accuracy but faster
 - Test specific code, no accuracy, fastest
- None of the open-source tools we've looked at dictate any specific methodology
 - The technology behind VProc has been used to add co-simulation capabilities to [OSVVM](#), giving access to a richer featured test environment (in VHDL), if working in that environment
- All of the tools mentioned don't rely on any of the others (mostly)
 - *mem_model* almost always goes with VProc, but it's not obligatory
 - *udplpPg* is based around VProc
 - the rv32 ISS can be compiled as an executable, and as a static library for integration into other environments such as C++ SoC models or SystemC etc.
- There are a rich set of open-source tools available, with some advanced capabilities, on which to develop you're project. We have looked at just a few today, integrated together to get something bigger than the sum of their parts.

Open Source Components and More Information



- The [Wireguard FPGA project](#) ([Chili.CHIPS*ba](#) Wireguard FPGA team)
 - Support from [NLnet Foundation](#)
- [mem_model](#) co-simulation sparse memory model ([Wyvern Semiconductors](#))
- [VProc](#) virtual processor co-simulation VIP (Wyvern Semiconductors)
 - [Article](#) on *VProc*, its use and how it functions
 - [PLI article](#) with *VProc* and *mem_model* as examples
- The [rv32](#) RISC-V instruction set simulator (Wyvern Semiconductors)
 - [Article](#) on *rv32* and its internal architecture
 - [Article](#) on *rv32*'s gdb remote debug interface
- [udpPglp](#) gigabit ethernet UDP/IPv4 RGMII/GMII co-simulation model (Wyvern Semiconductors)
 - Based on [tcpIpPg](#) (Wyvern Semiconductors)
- [peakrdl](#) and [systemrdl-compiler](#) (Alex Mykyta et. al)
- [Verilator](#) Verilog and SystemVerilog cycle based logic simulator (Wilson Snyder et. al)
- [GTKwave](#) (Anthony Bybell) and [Surfer](#) (Linköping University) waveform viewers

Wireguard FPGA Project Co-Simulation



END

Backup Slides

peakrdl c-header CSR access output



- Provides the macros for masking and shifting
- Provides a bitfield structure of the fields within a register
- Union for structure or whole register access
- Provides hierarchical structures (not shown)

```
#ifdef __cplusplus
extern "C" {
#endif

#include <stdint.h>
#include <assert.h>

// Reg - csr::ip_lookup_engine::table::ip_address
#define CSR_IP_LOOKUP_ENGINE_TABLE_IP_ADDRESS_ADDRESS_bm 0xffffffff
#define CSR_IP_LOOKUP_ENGINE_TABLE_IP_ADDRESS_ADDRESS_bp 0
#define CSR_IP_LOOKUP_ENGINE_TABLE_IP_ADDRESS_ADDRESS_bw 32
#define CSR_IP_LOOKUP_ENGINE_TABLE_IP_ADDRESS_MASK_bm 0xffffffff00000000
#define CSR_IP_LOOKUP_ENGINE_TABLE_IP_ADDRESS_MASK_bp 32
#define CSR_IP_LOOKUP_ENGINE_TABLE_IP_ADDRESS_MASK_bw 32
typedef union {
    struct __attribute__((__packed__)) {
        uint64_t address :32;
        uint64_t mask :32;
    } f;
    uint64_t w;
} csr_ip_lookup_engine_table_ip_address_t;

// Reg - csr::ip_lookup_engine::table::public_key
#define CSR_IP_LOOKUP_ENGINE_TABLE_PUBLIC_KEY_KEY_bm 0xffffffff
#define CSR_IP_LOOKUP_ENGINE_TABLE_PUBLIC_KEY_KEY_bp 0
#define CSR_IP_LOOKUP_ENGINE_TABLE_PUBLIC_KEY_KEY_bw 32
typedef union {
    struct __attribute__((__packed__)) {
        uint32_t key :32;
    } f;
    uint32_t w;
} csr_ip_lookup_engine_table_public_key_t;
```

Co-simulation HAL for Target (RISC-V)



- Python script written to generate common HAL using *systemrdl-compiler*
- For target, wraps *peakrdl* output in simple inline function calls
 - Hierarchical dereferencing
 - Function name matches field
 - Can access 'full' register
- Why not use *peakrdl* output directly?
 - We need a common API for native vs target
 - Bitfields can't be used in 'classic' method of overloading accesses to a given type to call *VProc* API functions

```
class csr_ip_lookup_engine_table_ip_address_vp_t {
public:
    csr_ip_lookup_engine_table_ip_address_vp_t (uint32_t* reg_addr = 0) :
        reg((csr_ip_lookup_engine_table_ip_address_t*)reg_addr) {};

    inline void      full(const uint64_t data) {reg->w = data;};
    inline uint64_t  full()                    {return reg->w;};

    inline void      address(const uint64_t data) {reg->f.address = data;};
    inline uint64_t  address()                    {return reg->f.address;};
    inline void      mask(const uint64_t data)    {reg->f.mask = data;};
    inline uint64_t  mask()                      {return reg->f.mask;};

private:
    csr_ip_lookup_engine_table_ip_address_t* reg;
};
```

```
#include "csr_hw.h"

int main (int argc, char** argv)
{
    // Create a CSR register object
    csr_vp_t* csr = new csr_vp_t(CSR_BASE_ADDR);

    csr->ip_lookup_engine->table[0]->allowed_ip[0]->address(0x12345678);
    uint32_t val = csr->ip_lookup_engine->table[0]->allowed_ip[0]->address();

    csr->ip_lookup_engine->table[0]->allowed_ip[0]->mask(0x65ef9078);
    val = csr->ip_lookup_engine->table[0]->allowed_ip[0]->mask();
}
```


Co-simulation HAL for VProc



- Wireguard's Python script can also generate a *VProc* HAL
- Produces the same functions to access the registers, fields and hierarchy
 - Usage is exactly the same
- Now *VProc* API calls are made to do the reads and writes
 - Uses the *peakrdl* generated macros for masks and shift of data for correct access alignment.
 - Will mimic write only access for aligned fields, doing read-modify-writes only for misaligned fields.
- Other considerations for native compilation of application include wrapping up delay functions
 - In native application make calls to `VTick()` to advance simulation time

```
class csr_ip_lookup_engine__table__ip_address_vp_t {
public:
    csr_ip_lookup_engine__table__ip_address_vp_t (uint32_t* reg_addr = 0) :
        reg((uint64_t)reg_addr) {};

    inline void address (const uint64_t data) {
        uint32_t wdata = (uint32_t)(data & 0xffffffff);
        VWriteBE(reg + 0, wdata << 0, 0xf, NO_DELTA_UPDATE, SOC_CPU_VPNODE);
        VTick(rand() % 33, SOC_CPU_VPNODE); };

    inline uint64_t address () {
        uint32_t rdata;
        VRead(reg + 0, &rdata, NO_DELTA_UPDATE, SOC_CPU_VPNODE);
        VTick(rand() % 33, SOC_CPU_VPNODE);
        return (((uint64_t)rdata << 0) &
            CSR_IP_LOOKUP_ENGINE__TABLE__IP_ADDRESS__ADDRESS_bm) >>
            CSR_IP_LOOKUP_ENGINE__TABLE__IP_ADDRESS__ADDRESS_bp; };

    inline void mask (const uint64_t data) {
        uint32_t wdata = (uint32_t)(data & 0xffffffff);
        VWriteBE(reg + 4, wdata << 0, 0xf, NO_DELTA_UPDATE, SOC_CPU_VPNODE);
        VTick(rand() % 33, SOC_CPU_VPNODE); };

    inline uint64_t mask () {
        uint32_t rdata;
        VRead(reg + 4, &rdata, NO_DELTA_UPDATE, SOC_CPU_VPNODE);
        VTick(rand() % 33, SOC_CPU_VPNODE);
        return (((uint64_t)rdata << 32) &
            CSR_IP_LOOKUP_ENGINE__TABLE__IP_ADDRESS__MASK_bm) >>
            CSR_IP_LOOKUP_ENGINE__TABLE__IP_ADDRESS__MASK_bp; };

private:
    uint32_t reg;
}
```